

# AIML427 Big Data: Assignment 3

Devon Grossett

2024-06-11

## Using Spark Machine Learning Libraries on Covertypes Dataset

### Introduction

For this project I have decided to use the Covertypes dataset from the UCI Machine Learning Repository. Before processing, this dataset contains 54 features and consists of approximately 500,000 observations. Each observation is a 30 x 30m cell of forest in the Roosevelt National Forest of Northern Colorado.

The dataset contains cartographic information such as the cell elevation, slope (in degrees), aspect (compass direction of slope in degrees), horizontal and vertical distance to surface water features (in meters), horizontal distance to nearest roadway, horizontal distance to nearest wildfire ignition point, and the hillshade index at 9am, noon, and 3pm at summer solstice, which is an index from 0 to 255 (with 0 being complete shade and 255 complete sun). All of these features are integers, and there are 10 of them in total.

As well as the cartographic variables, there are two categorical variables which have been one-hot encoded. The first is `Wilderness_Area`, which has four levels representing different wilderness areas (Rawah, Neota, Comanche Peak, and Cache la Poudre). The other is `Soil_Type` which has 40 levels representing different soil families. Each soil family has a four digit USFS (United States Forestry Service) Ecological Landtype Unit code. The first digit of this code represents the climatic zone of the soil, and the second digit represents the geological zone. In total, there are 44 of the one-hot encoded variables derived from these two categorical variables and they are all integer valued. Combined with the 10 numeric features, this gives a total of 54 original features.

Lastly there is the target feature we are trying to predict the value of, which is a categorical variable representing the forest cover type. It can take on one of seven different values such as 'Spruce/Fir', 'Lodgepole Pine', 'Cottonwood/Willow', etc. It is coded as an integer taking values from 1 to 7.

All of this information can be found in the `covtype.info` file that comes with the downloaded data, and the downloaded data is 75.2MB in size before any preprocessing.

In order to perform the task of predicting the forest cover type, we will be using the Python API for Apache Spark, PySpark. There are several options for Multiclass classification in PySpark, but I have decided to use the `LogisticRegression` class of model from the

`pyspark.ml.classification` library. This class of model supports multiclass classification through the use of a Multinomial Logit model (a Generalised Linear Model with a Multinomial distribution family, and a Softmax link function). For each observation, this model calculates a probability for each class, and classifies the observation belonging to the class with the largest predicted probability.

The distribution of class labels is shown in Figure 1. Here we can see that the vast majority of our data falls into the first two class labels, having over 200,000 observations each. The remaining all have 50,000 or fewer observations.




Figure 1: Distribution of class labels

## Preprocessing

The first step in preprocessing was to select reference levels for the one-hot encoded `Wilderness_Area` and `Soil_Type` variables. Without this step, the design matrix will be non-invertible and any attempt to find coefficients will have trouble converging. For this analysis we just chose the first variable of each (`Wilderness_Area_1` and `Soil_Type_1`) to be the reference level and removed them from our dataset. This choice does not affect the model, just how we interpret the coefficients.

We have also constructed the features `cos_aspect` and `sin_aspect` by converting `Aspect` to radians and taking the sine and cosine of this angle. This is due this feature representing the compass direction that the slope faces, so an aspect of  $0^\circ$  and  $360^\circ$  both represent exactly

north. By taking the cosine,  $0^\circ$  and  $360^\circ$  both transform to 1.0, and `cos_aspect` represents “Northness” and avoids the issue of discontinuity highlighted above. Likewise for `sin_aspect`, except now 1 represents East and -1 West.

Several interaction features were also constructed: `Slope` was interacted with `Elevation`, `cos_aspect`, and `sin_aspect`, and each of the `Hillshade` features were interacted with each other.

After these steps were completed, there were 60 features in total and the data was larger than 20MB in size.

Additionally, we wished to compare the performance of the Multinomial Logistic Regression model with and without standardizing the features to have a mean of zero and a standard deviation of one. So an additional preprocessing step of standardizing the numeric features was done. The binary features weren’t standardized, as they already fall in the range  $[0, 1]$ , and by standardizing we lose the useful interpretation of the coefficients being the average difference in the linear predictor for that feature versus its reference level.

The final preprocessing was to reduce the original set of features down to a smaller number of principal components, to compare the model performance before and after performing this step. From Figure 2, we can see that by using the first 10 PC’s of our data, we capture 94% of the variance of the original features. Hence this was chosen to be the number of principal components we used.




Figure 2: Variance of original data explained by first 10 principal components

## Readme.txt

```
## Move Files to ECS System
```

```
Move multinomial_logistic_py, covtype.data, and SetupSparkClasspath.sh to
barretts@ecs.vuw.ac.nz
```

```
$ scp multinomial_logistic.py <username>@barretts.ecs.vuw.ac.nz:
etc..
```

### ## Access VUW Hadoop cluster

```
ssh into barretts using your ecs account
$ ssh <username>@barretts.ecs.vuw.ac.nz
```

```
ssh into one of the Hadoop nodes
$ ssh co246a-5
(last number can be 1-8)
```

### ## Setup Hadoop and Spark

```
configure Hadoop and Spark
$ source SetupSparkClasspath.sh
```

```
create directory for input and output datasets
$ hadoop fs -mkdir /user/<username>/input /user/<username>/output
```

```
upload input data into hdfs
$ hadoop fs -put covtype.data /user/<username>/input/
```

### ## Run Spark Job

multinomial\_logistic.py takes 3 inputs:

- path to input data
- path to output folder
- random seed

```
$ spark-submit --master yarn --deploy-mode cluster multinomial_logistic.py
/user/<username>/input/covtype.data /user/grossedevo/output 200
```

### ## Retrieve Results

```
move from hdfs to ecs local
$ hadoop fs -copyToLocal /user/<username>/output
$ hadoop fs -rm -r /user/<username>/output
```

```
move from ECS system to desired path local pc
$ scp -r <username>@barretts.ecs.vuw.ac.nz:~/output ~/path/to/local
```

## Modelling Pseudo-code

```
if not enough arguments:
    exit

initialize spark session

data = read(csv file, column names)
drop = drop(reference levels of categorical features) from data

create unscaled features

test_data, train_data = 80:20 test train split of data

fit standard scaler on train data
transform test_data and train_data with standard scaler

fit pca model on train data
transform test data and train data with pca model

initialize F_1 evaluator

for each set of features:
    record start time
    fit logistic regression to unstandardized train_data
    record end time
    train_time = start_time - end_time

    train_preds = transform train_data with lr_model
    test_preds = transform test_data with lr_model

    train_f1 = evaluator(train_preds)
    test_f1 = evaluator(test_preds)

for each set of features:
    store train_f1, test_f1, train_time in dataframe
write results dataframe to storage

for each set of features:
    store model coefficients in dataframe
write coefficients dataframe to disk

store confusion matrices in dataframe
write confusion matrices dataframe to disk

for pca model:
    store principal component weights in dataframe
write principal component weights dataframe to disk
```

```

store percent of variance explained to dataframe
write percent of variance explained dataframe to disk

stop spark session

```

## Comparing Unscaled and Scaled Model

Prior to finalising the model to be used for training, I performed some tests to determine if regularization would be beneficial for this problem. This was done via a grid search over `regParam` and `elasticNetParam` ( $\lambda$  and  $\alpha$ , respectively). The results are shown in Figure 3, and suggest that the best performance is with either a very small  $\lambda$ , or  $\lambda = 0$ . For simplicity, we will use  $\lambda = 0$  i.e. no regularization.

Sum - F1 Score	elasticNetParam				
regParam	0	0.25	0.5	0.75	1
0.000E+00	0.7137	0.7137	0.7137	0.7137	0.7137
1.000E-05	0.7137	0.7125	0.7131	0.7131	0.7130
1.585E-05	0.7135	0.7128	0.7127	0.7126	0.7122
2.512E-05	0.7133	0.7127	0.7126	0.7126	0.7127
3.981E-05	0.7133	0.7127	0.7123	0.7123	0.7124
6.310E-05	0.7133	0.7125	0.7121	0.7120	0.7125
1.000E-04	0.7126	0.7122	0.7120	0.7119	0.7121
1.585E-04	0.7130	0.7122	0.7125	0.7127	0.7120
2.512E-04	0.7120	0.7122	0.7121	0.7120	0.7124
3.981E-04	0.7112	0.7129	0.7124	0.7123	0.7130
6.310E-04	0.7103	0.7103	0.7117	0.7121	0.7124
1.000E-03	0.7088	0.7092	0.7098	0.7111	0.7119
1.585E-03	0.7067	0.7075	0.7089	0.7095	0.7094
2.512E-03	0.7046	0.7055	0.7066	0.7073	0.7072
3.981E-03	0.7023	0.7035	0.7038	0.7040	0.7031
6.310E-03	0.7001	0.6999	0.7001	0.7008	0.6976
1.000E-02	0.6980	0.6955	0.6914	0.6867	0.6818
1.585E-02	0.6941	0.6856	0.6787	0.6740	0.6705
2.512E-02	0.6863	0.6725	0.6708	0.6660	0.6619
3.981E-02	0.6774	0.6663	0.6581	0.6527	0.6472
6.310E-02	0.6659	0.6474	0.6443	0.6369	0.6106
1.000E-01	0.6532	0.6297	0.6116	0.5642	0.5614

Figure 3: Validation  $F_1$  scores for grid search. Suggests no regularization is needed

For model evaluation, we use the test  $F_1$  score, which balances recall and precision by taking their harmonic mean. Our data is split into 80% training, and 20% test.

With the model form now finalised, we compare the results with and without standardization of the features. Table 1 provides the training and test scores, as well the run time (only training the model, does not include additional preprocessing for the standardized model). As can be seen, both models achieved identical training and test scores (to 3 decimal places), with the only difference being a minor difference in execution time.

Table 1: Training and Test  $F_1$  scores for the Standardized and Unstandardized models

Measure	Value	Value (Standardized)
Training $F_1$	0.717	0.717
Test $F_1$	0.720	0.720
Run time (m)	2.167	2.201

For the random seed used for this run (200), we obtained slightly higher test scores than

training scores, although they are very close. This indicates that in both cases there is no overfitting happening and our model generalises well to unseen data.

The training times for the two models were also almost identical, so when taking into account the additional time to fit the **StandardScaler** and transform the training and test set the Standardized model was more expensive to fit.

From the confusion matrices (see Appendix), we are able to calculate the precision and recall scores for each label in our test set. These are given in Table 2. It appears that labels 4, 5, and 6 (Cottonwood/Willow, Aspen, and Douglas-fir) are the most problematic for our model. Aspen in particular has a recall of only 3%, meaning that the model very rarely predicts this Cover Type even when it is the true class label.

Table 2: Per-Label Precision and Recall for the Unstandardized Model

Class	Precision	Recall
1	0.713	0.702
2	0.752	0.802
3	0.697	0.796
4	0.658	0.444
5	0.504	0.033
6	0.477	0.283
7	0.741	0.582

The models obtained by these two methods give very similar predictions, but have different interpretation of their coefficients. The models are quite large, as there are 7 equations (one for each class label), with 61 coefficients each. Hence I have included them in the appendix.

For the unstandardized model, we can interpret the coefficients as the increase to the linear predictor from a one-unit increase in the independent variable. For the standardized model, the interpretation of the coefficients changes to the increase to the linear predictor from an increase of one standard deviation increase in the independent variable.

## Comparing Unscaled and PCA Model

In order to compare the model with and without applying PCA to the features, a pyspark PCA model was fit against the training data (to ensure that there was no data leakage), and then both the training and test sets were transformed with this model.

Table 3 gives the results for our original model again, and the results for the PCA model. There is a slight reduction in test scores for the PCA transformed model, but a significant speed up in training time.

Table 3: Training and Test  $F_1$  scores for the Unstandardized and PCA models

Measure	Value	Value (PCA)
Training $F_1$	0.717	0.690
Test $F_1$	0.720	0.692
Run time (m)	2.167	0.450

Compared to the unstandardized model, the PCA transformed model experienced a 3.9% decrease in its test  $F_1$  score. This decrease in test score came with the benefit of only taking 21% of the time to train as the unstandardized model. This speed-up is only for the time to train the model, and did not take into account the time to fit the PCA transformer and transform the datasets. However, this is a one time cost, and if there was a need to train the model on data regularly this time save would begin to quickly outweigh the initial compute cost.

The model coefficients for the PCA transformed model are given in the appendix.

Compared to our original model, the PCA transformed model has only 10 coefficients + the intercept. This number was selected by us, but it is a much simpler model than the original while only sacrificing a small amount of accuracy. However, the model does become less interpretable, as it is a lot harder to explain variables in terms of principal components versus their original forms. This is the trade-off that comes with the simpler model.

Table 4 gives the per-label precision and recall scores for our PCA transformed model. Again we see deficiencies in scores for labels 4, 5, and 6. Aspen

(label 5) is particularly problematic again, this time being unable to identify any true positives, leading to precision and recall scores of zero. If we had a particular interest in being able to correctly identify this cover type, then we may want to go back to the drawing board to work out why this label is proving difficult for our model to predict.

Table 4: Per-Label Precision and Recall for the PCA Model

Class	Precision	Recall
1	0.684	0.697
2	0.736	0.788
3	0.656	0.737
4	0.671	0.256
5	0.000	0.000
6	0.407	0.217
7	0.642	0.339

## Conclusion

In conclusion, we saw no discernable difference in model performance between the data trained on unstandardized and standardized data, but the standardizing required an additional preprocessing step. On the other hand, reducing the dimensionality of the data using the first ten principal components saw a slight reduction in test performance but a significant speed up in the time taken to train the model. This trade-off may become worth considering if our model was more complex, or had more data and took considerable time to train initially. However in this case, the model was not complex or the data overly large so the benefit was less noticable, and if a model needed to be selected the the first model on unstandardized data would be my choice.

When examining the class specific performance, we saw a decrease in test precision and recall for Aspen, Douglas-fir, and Krummholz. This could potentially something that would need addressing in any further follow up modelling depending on the use case for such a model. For the meantime we just note it as a deficiency in the model predictions.



## Appendix

Table 5: Unstandardized model confusion matrix (test data)

Actual \ Predicted	1	2	3	4	5	6	7
1	29,847	11,834	13	0	3	14	793
2	10,312	45,624	542	1	32	326	38
3	0	686	5,668	111	16	636	0
4	1	1	231	244	0	74	0
5	12	1,725	41	0	61	17	0
6	0	800	1,641	15	9	975	0
7	1,674	31	0	0	0	0	2,372

Table 6: Standardized model confusion matrix (test data)

Actual \ Predicted	1	2	3	4	5	6	7
1	29,850	11,834	13	0	3	15	789
2	10,316	45,622	539	1	34	326	37
3	0	686	5,667	111	16	637	0
4	0	1	233	242	0	74	0
5	12	1,723	41	0	63	17	0
6	0	799	1,642	14	8	977	0
7	1,672	31	0	0	0	0	2,374

Table 7: PCA model confusion matrix (test data)

Actual \ Predicted	1	2	3	4	5	6	7
1	29,629	12,084	12	0	0	16	763
2	11,072	44,810	634	0	9	343	7
3	0	1,194	5,247	48	0	628	0
4	0	11	316	141	0	82	0
5	2	1,817	17	0	0	20	0
6	0	903	1,769	21	0	747	0
7	2,622	73	0	0	0	0	1,382

Table 8: Coefficients for Unstandardized Model

Variable	Class 1	Class 2	Class 3	Class 4	Class 5	Class 6	Class 7
Intercept	- 12.158	6.687	23.483	35.929	-1.595	7.284	- 59.629
Elevation	0.010	0.003	-0.011	-0.020	-0.001	-0.011	0.029

Variable	Class 1	Class 2	Class 3	Class 4	Class 5	Class 6	Class 7
Slope	-0.343	-0.043	0.238	-0.155	-0.264	0.124	0.443
Elevation:Slope	0.000	0.000	0.000	0.000	0.000	0.000	0.000
cos_aspect	0.213	0.119	-0.904	0.030	0.579	-0.396	0.360
sin_aspect	0.090	0.067	-0.246	0.500	-0.759	-0.027	0.375
Slope:cos_aspect	-0.008	-0.016	0.006	-0.075	0.062	0.100	-0.069
Slope:sin_aspect	-0.094	-0.097	0.001	0.067	0.025	0.032	0.066
Horizontal_Dist_Hydrology	-0.001	0.001	0.004	-0.001	0.000	0.001	-0.003
Vertical_Dist_Hydrology	-0.004	-0.003	0.001	0.009	0.002	-0.001	-0.005
Horizontal_Dist_Roadways	0.000	0.000	0.000	0.002	-0.001	0.000	0.000
Horizontal_Dist_Fire_Points	0.000	0.000	0.000	0.001	0.000	0.000	0.000
Hillshade_9am	0.008	-0.013	0.010	0.008	0.030	0.058	-0.102
Hillshade_Noon	-0.041	-0.026	0.022	0.025	-0.004	0.048	-0.023
Hillshade_3pm	-0.014	-0.045	0.082	-0.006	-0.066	0.064	-0.014
Hillshade_9am:Hillshade_Noon	0.000	0.000	0.000	0.000	0.000	0.000	0.000
Hillshade_9am:Hillshade_3pm	0.000	0.000	0.000	0.000	0.000	0.000	0.000
Hillshade_Noon:Hillshade_3pm	0.000	0.000	0.000	0.000	0.000	0.000	0.000
Wilderness_Area_2	-1.032	-0.145	1.020	7.568	-6.929	0.453	-0.934
Wilderness_Area_3	-3.001	-3.104	3.735	-0.755	-2.514	5.929	-0.290
Wilderness_Area_4	-9.055	-2.387	3.788	11.724	-8.859	5.623	-0.834
Soil_Type_2	-	4.524	2.363	1.930	4.126	1.159	-3.280
	10.822						
Soil_Type_3	-	7.923	4.600	6.103	-9.059	2.969	-2.519
	10.016						
Soil_Type_4	-2.156	0.741	-0.708	-0.244	0.434	-2.808	4.743
Soil_Type_5	-7.940	-	7.999	8.421	-4.347	8.265	-0.453
		11.945					
Soil_Type_6	-8.613	6.617	3.301	3.583	-6.090	2.373	-1.171
Soil_Type_7	-	9.069	2.191	3.130	-3.313	0.971	-0.957
	11.090						
Soil_Type_8	3.004	5.662	-0.260	0.727	-4.329	-1.185	-3.619
Soil_Type_9	4.854	6.192	-4.327	5.764	-7.028	-4.088	-1.366
Soil_Type_10	0.766	3.650	-0.730	-0.254	1.276	-1.146	-3.562
Soil_Type_11	0.827	4.124	-0.383	0.685	2.689	-2.035	-5.906
Soil_Type_12	3.766	7.257	-1.901	0.482	-5.518	-1.724	-2.361
Soil_Type_13	-0.847	2.535	-4.043	4.382	1.658	-1.344	-2.341
Soil_Type_14	-9.208	-	10.037	13.157	-6.136	10.618	-1.234
		17.233					
Soil_Type_15	-7.908	-	-1.994	7.403	-3.274	17.027	-0.297
		10.957					
Soil_Type_16	0.993	3.373	-0.877	1.207	1.439	-0.900	-5.233
Soil_Type_17	0.212	2.751	-0.420	2.460	3.483	-0.874	-7.611
Soil_Type_18	3.021	5.989	-4.907	-1.776	5.240	-5.233	-2.333
Soil_Type_19	1.941	4.324	-5.389	0.568	4.650	-6.482	0.388
Soil_Type_20	2.540	4.736	-4.436	-0.024	2.163	-0.266	-4.713
Soil_Type_21	3.345	2.645	-4.376	7.352	-5.802	-5.637	2.473

Variable	Class 1	Class 2	Class 3	Class 4	Class 5	Class 6	Class 7
Soil_Type_22	2.303	4.394	-3.478	6.484	-5.854	-5.420	1.571
Soil_Type_23	0.649	3.000	-6.537	2.270	2.077	-2.332	0.873
Soil_Type_24	1.279	4.138	-8.041	1.019	1.649	-1.185	1.142
Soil_Type_25	-0.272	3.985	0.969	7.153	-4.800	1.277	-8.312
Soil_Type_26	3.092	6.690	-6.762	4.899	6.114	-7.206	-6.827
Soil_Type_27	1.831	4.689	-3.880	6.892	-6.970	-4.589	2.027
Soil_Type_28	3.047	6.529	-9.431	6.946	4.384	-9.480	-1.994
Soil_Type_29	0.276	3.395	-3.426	-1.196	1.955	-3.533	2.529
Soil_Type_30	0.372	3.545	-4.159	-2.860	3.298	-3.192	2.996
Soil_Type_31	0.403	3.453	-8.864	6.688	1.704	-3.125	-0.258
Soil_Type_32	-1.389	2.098	-2.543	6.020	0.231	-2.398	-2.019
Soil_Type_33	-0.286	2.749	-6.495	6.690	0.572	-2.531	-0.698
Soil_Type_34	-1.576	3.403	-9.868	6.438	1.101	-0.616	1.117
Soil_Type_35	-1.064	-0.930	-0.407	7.413	-5.136	-0.902	1.026
Soil_Type_36	-1.737	2.773	-2.797	7.205	-5.975	-1.029	1.560
Soil_Type_37	-	-	2.716	7.592	-3.387	1.571	18.444
	14.636	12.300					
Soil_Type_38	-0.919	1.039	-2.460	7.382	-5.167	-1.497	1.622
Soil_Type_39	0.401	1.905	-2.904	7.446	-5.308	-4.779	3.240
Soil_Type_40	-1.379	0.704	-0.799	7.413	-5.120	-1.402	0.583

Table 9: Coefficients for Standardized Model

Variable	Class 1	Class 2	Class 3	Class 4	Class 5	Class 6	Class 7
Intercept	8.876	6.881	-1.990	-	3.643	-3.291	-0.452
				13.666			
Elevation	2.927	0.944	-2.985	-5.484	-0.348	-3.074	8.020
Slope	-2.538	-0.243	1.740	-1.033	-1.802	0.889	2.988
Elevation:Slope	2.396	0.388	-2.148	0.484	2.521	-0.828	-2.812
cos_aspect	0.144	0.081	-0.621	0.030	0.393	-0.265	0.237
sin_aspect	0.060	0.045	-0.170	0.337	-0.507	-0.018	0.253
Slope:cos_aspect	-0.096	-0.185	0.095	-0.802	0.672	1.123	-0.807
Slope:sin_aspect	-0.972	-1.001	-0.054	0.653	0.379	0.301	0.695
Horizontal_Dist_Hydrology	-0.277	0.125	0.854	-0.220	-0.077	0.273	-0.678
Vertical_Dist_Hydrology	-0.221	-0.153	0.062	0.527	0.107	-0.030	-0.293
Horizontal_Dist_Roadways	-0.552	-0.411	0.002	2.647	-1.142	0.165	-0.710
Horizontal_Dist_Fire_Points	0.320	-0.366	-0.627	1.699	-0.547	0.027	0.133
Hillshade_9am	0.237	-0.360	0.200	0.218	0.861	1.449	-2.606
Hillshade_Noon	-0.807	-0.493	0.416	0.494	-0.033	0.934	-0.512
Hillshade_3pm	-0.477	-1.667	2.979	-0.237	-2.340	2.256	-0.514
Hillshade_9am:Hillshade_Noon	0.044	0.147	0.132	-0.543	0.452	-0.736	0.593
Hillshade_9am:Hillshade_3pm	0.293	0.811	-0.815	-0.349	-0.081	-0.611	0.753
Hillshade_Noon:Hillshade_3pm	0.107	0.348	-2.519	0.699	3.306	-0.801	-1.139

Variable	Class 1	Class 2	Class 3	Class 4	Class 5	Class 6	Class 7
Wilderness_Area_2	-1.106	-0.221	0.980	7.806	-6.896	0.442	-1.003
Wilderness_Area_3	-2.974	-3.080	3.849	-1.100	-2.487	6.047	-0.256
Wilderness_Area_4	-9.124	-2.492	3.764	11.963	-8.849	5.598	-0.859
Soil_Type_2	-	4.564	2.357	1.939	4.211	1.161	-3.323
	10.909						
Soil_Type_3	-	8.004	4.645	6.155	-9.177	3.026	-2.555
	10.100						
Soil_Type_4	-2.125	0.717	-0.766	-0.279	0.447	-2.853	4.859
Soil_Type_5	-8.011	-	8.086	8.497	-4.427	8.366	-0.475
		12.035					
Soil_Type_6	-8.688	6.699	3.337	3.612	-6.185	2.423	-1.198
Soil_Type_7	-	9.401	2.163	3.043	-3.387	0.941	-0.984
	11.179						
Soil_Type_8	3.139	5.760	-0.292	0.692	-4.406	-1.228	-3.666
Soil_Type_9	4.890	6.190	-4.390	5.954	-7.122	-4.129	-1.394
Soil_Type_10	0.814	3.658	-0.757	-0.277	1.326	-1.159	-3.606
Soil_Type_11	0.889	4.145	-0.404	0.630	2.750	-2.040	-5.970
Soil_Type_12	3.831	7.285	-1.976	0.646	-5.612	-1.775	-2.400
Soil_Type_13	-0.891	2.448	-4.183	4.707	1.609	-1.452	-2.239
Soil_Type_14	-9.285	-	10.134	13.263	-6.230	10.727	-1.260
		17.349					
Soil_Type_15	-7.979	-	-2.038	7.626	-3.347	17.096	-0.318
		11.040					
Soil_Type_16	1.040	3.382	-0.921	1.173	1.531	-0.917	-5.288
Soil_Type_17	0.266	2.767	-0.446	2.451	3.540	-0.892	-7.686
Soil_Type_18	3.236	6.177	-4.991	-2.183	5.453	-5.321	-2.371
Soil_Type_19	1.902	4.243	-5.380	0.652	4.617	-6.446	0.412
Soil_Type_20	2.577	4.734	-4.449	-0.019	2.213	-0.288	-4.768
Soil_Type_21	3.295	2.562	-4.393	7.576	-5.885	-5.548	2.393
Soil_Type_22	2.212	4.262	-3.392	6.684	-5.926	-5.316	1.475
Soil_Type_23	0.656	2.967	-6.507	2.352	2.083	-2.438	0.887
Soil_Type_24	1.288	4.108	-8.235	1.200	1.685	-1.230	1.183
Soil_Type_25	-0.262	3.961	0.933	7.378	-4.883	1.273	-8.399
Soil_Type_26	2.995	6.548	-6.697	5.169	6.011	-7.127	-6.900
Soil_Type_27	1.692	4.515	-3.752	7.118	-7.057	-4.411	1.895
Soil_Type_28	3.070	6.498	-9.476	7.165	4.365	-9.597	-2.026
Soil_Type_29	0.340	3.417	-3.501	-1.266	2.016	-3.615	2.609
Soil_Type_30	0.418	3.548	-4.208	-2.945	3.338	-3.228	3.077
Soil_Type_31	0.396	3.407	-8.964	6.919	1.690	-3.185	-0.262
Soil_Type_32	-1.423	2.025	-2.617	6.364	0.204	-2.502	-2.050
Soil_Type_33	-0.306	2.691	-6.573	6.950	0.556	-2.614	-0.706
Soil_Type_34	-1.612	3.328	-9.833	6.641	1.086	-0.692	1.083
Soil_Type_35	-1.185	-1.135	-0.306	7.641	-5.135	-0.797	0.918
Soil_Type_36	-1.881	2.602	-2.671	7.431	-6.007	-0.912	1.438
Soil_Type_37	-	-	2.691	7.821	-3.461	1.570	18.502
	14.735	12.388					

Variable	Class 1	Class 2	Class 3	Class 4	Class 5	Class 6	Class 7
Soil_Type_38	-1.075	0.843	-2.277	7.612	-5.202	-1.379	1.479
Soil_Type_39	0.341	1.802	-2.770	7.678	-5.371	-4.870	3.192
Soil_Type_40	-1.510	0.518	-0.702	7.642	-5.140	-1.264	0.456

Table 10: Coefficients for PCA Model

Variable	Class 1	Class 2	Class 3	Class 4	Class 5	Class 6	Class 7
Intercept	6.0061	6.9643	-2.9736	-8.0154	2.6069	-2.5043	-2.0839
PC1	-0.1453	-0.0892	0.1116	0.2715	0.1167	0.0142	-0.2796
PC2	-0.1877	-0.0071	0.5720	-0.1821	0.0750	0.6607	-0.9307
PC3	0.2214	0.0051	-0.6876	0.5455	-0.2540	0.0379	0.1317
PC4	1.5511	0.6194	-1.8136	-2.2839	-0.1940	-1.6386	3.7597
PC5	-0.6894	-0.1832	1.5211	-0.6586	0.4410	1.0568	-1.4877
PC6	-2.4725	-0.7590	2.6060	4.6188	-0.3661	2.2723	-5.8997
PC7	2.0547	0.9262	-2.0286	-4.6901	1.0392	-1.6206	4.3192
PC8	-0.5748	0.1544	1.3865	-0.0529	-0.2167	0.8234	-1.5199
PC9	1.2178	0.9703	-2.3286	1.5692	0.5496	-2.4313	0.4530
PC10	-0.4170	-0.2494	-0.3349	1.0290	0.6848	0.1034	-0.8158